

Beer Preferences for Thursday Night Football: A Balanced Incomplete Block Design Experiment

Matthew Murnane

Ethan Newcomb

Abstract

A Balanced Incomplete Block Design was utilized during a weekly Thursday night football watch party to test a friend groups beer taste. The study involved four participants and four kinds of beer. The H_0 stated that all beers were liked equally and H_1 stated that at least one beer was not liked equally. Ultimately, we failed to reject the null hypothesis with high p-values. Suprisingly, there was no block effect. The data collected in this experiment could be utilized as a helpful prior in variance estimation and power calculations in the future.

Introduction

American football has immensely grown in popularity in recent years. Along with the growth of a national sport comes an increase in social gatherings and alcoholic beverage enjoyment. Weekly Thursday night football games are broadcast on national television and each Thursday, Matthew's house hosts a watch party along with his roommates. The group enjoys consuming a wide variety of beer, ranging in type and brand. This setting became ideal to test the following research question: What is the beer type that is most favored among the group of friends? Since each individual holds their own inherent biases toward beer, a Balanced Incomplete Block Design (BIBD) is a perfect experimental design. Here, the participants' preferences for beer are blocked, allowing the effects of beer type to be highlighted without the unwanted effect of the nuisance factor. The data collected will be the participants' ratings of beers, which differ in type and flavor. The "Balanced Incomplete" part of the experiment comes from the fact that alcohol impairs judgement, so the amount given to the participants should be limited.

Methods

There are four participants in this study: Zach, Jon, Nolan and Beni. All are male, in their young twenties and beer drinkers. These four participants constituted the entire population of interest in our study. Because of this there is no need for a mixed effects model where participants are random.

The constraints of a BIBD design are that when not all treatments can be assigned to each block a subset of those treatments are assigned. To ensure equal precision of estimation of each treatment: Each treatment appears in the same number of blocks, each treatment has the same replication and there are the same number of pairwise comparisons. A BIBD experiment satisfies the following relation:

$$a \cdot r = b \cdot k$$

where, a = number of treatments in the experiment,

r = number of blocks which any one treatment appears,

b = number of blocks in the experiment,

k = number of treatments per block

We are interested in testing four kinds of beers: A dark beer, pale ale, IPA, and a generic beer. The brands used for this were Guinness for the dark beer, Sierra Nevada for pale ale option, Pliny the Elder for IPA and Coors light for the generic beer. That is $a = 4$ treatment options. Our introduction explained the desire to limit the number of beers tasted because alcoholic consumption impairs judgement. To mediate the affect of alcohol on judgement we will limit the number of beers to $k = 3$. With the number of participants being $b = 4$ we will insure that any given beer appears only $r = 3$ across all blocks. We then have:

$$a \cdot r = b \cdot k$$

$$4 \cdot 3 = 4 \cdot 3$$

$$12 = 12$$

Satisfying the relation proved, we have a valid BIBD experiment.

Everything was randomized before starting the experiment. With simple R code (can be found in the Appendix) and using `set.seed(530)` we were able to get a correct BIBD set up found in Table 1.

Table 1: Blanced Incomplete Block Design Setup

Participant	I	II	III
1	C	B	D
2	B	C	A
3	D	A	C
4	D	B	A

We then randomized the assignments of beer types to the letters and the order of participants, using the same `set.seed(530)`. The table below is of treatment assignments and order assignments.

Table 2: Randomization of Treatments and Participants

Treatment	Beer	Order	Participant
A	Sierra	1	Zach
B	Coors	2	Jon
C	Guinness	3	Nolan
D	Pliny	4	Beni

The Materials used for the experiment were red solo cups, a pitcher of water, a tablet for recording the ratings, and a blindfold. The red solo cups allowed for the participants to not feel the can/bottle that the beer came in. This was necessary to keep them from making guesses that could affect their judgement.

The experiment was conducted inside of a separate room from the watch party. Participants were blind folded before entry into the room. The blindfold was necessary so that participants did not see what beers were being offered and what beers had already been open. It was best for the experiment that the only sense allowed was the sense of taste. They were led into the room, seated down and told the following:

You will be offered three beers during the experiment. You will be given a glass of water to cleanse your palate before the tasting begins and another glass of water in between each tasting of beer. After tasting you will rate the beer on a scale from 1 to 10. 1 meaning "I never want to drink this again", 5 meaning "this is an okay beer", and 10 meaning "I want a whole glass of this beer right now". Half points are allowed.

Participants were then given a glass of water to start and then a cup of beer in the kind and order dictated by Tables 1 & 2. After each tasting the participant was given as much time as needed to give the tasted beer an appropriate rating. Ratings were documented on the tablet device and transferred into the code found in the Appendix directly after the experiment concluded.

The Design

In summary, we conducted a Balanced Incomplete Block Design where the experimental units are the participants during a specific tasting, the measurement is the rating given by the participant after the tasting, the treatments are the four different kinds of beer and the block are the participants. A Balanced Incomplete Block Design in analyzed using ANOVA.

The hypothesis:

$$H_0 : \text{All kinds of beer are like equally, } \mu_i = \mu_{i'}, \forall i \neq i'$$

$$H_1 : \text{Atleast one beer is not liked equally, } \mu_i \neq \mu_{i'}, \text{ for atleast one pair } i \neq i'$$

The model:

$$y_{ij} = \mu + \tau_i + \beta_j + \epsilon_{ij} \quad \text{where } i = 1, 2, 3, 4 \text{ and } j = 1, 2, 3, 4$$

$$\text{and where, } \sum_{i=1}^4 \tau_i = 0 \text{ and } \sum_{j=1}^4 \beta_j = 0$$

Note that not all y_{ij} exist and we assume $\epsilon_{ij} \sim N(0, \sigma^2)$

Results

The ratings of each participant's three randomly assigned beers are shown in Table 3. Additionally, the mean rating of each beer type and participant are displayed. It can be observed that the beer type with the highest mean rating is Coors (Generic) at 6.17. On the other hand, Guinness (Dark) received the lowest mean rating at 4.33. The range rating between the highest and lowest mean ratings is 1.84, indicating a relatively small difference between mean ratings of beer type. Regarding participant ratings, the highest mean rating resulted from Beni at 6.5, while the lowest mean rating resulted from Zach at 3.67. The range between these two mean ratings is 2.83, indicating a relatively small difference in participant mean ratings.

There appear to be apparent outliers, such as the ratings of 8.5 or 2, which could indicate a significant block effect. However, the ANOVA table results at the end of this section will disprove that notion.

Table 3: Beer Ratings by Participant

Beer	Nolan	Jon	Beni	Zach	Row Means
Sierra(Pale Ale)	NA	5.5	7	2.5	5
Coors(Generic)	6	6	NA	6.5	6.17
Guinness(Dark)	5.5	3.5	4	NA	4.33
Pliny(IPA)	3	NA	8.5	2	4.5
Column Means	4.83	5	6.5	3.67	$\hat{\mu} = 5$

Figure 1 displays a visual representation of Ratings Per Beer. It is relatively difficult to draw a conclusion from this data, given the wide range of data points and their overlapping nature.

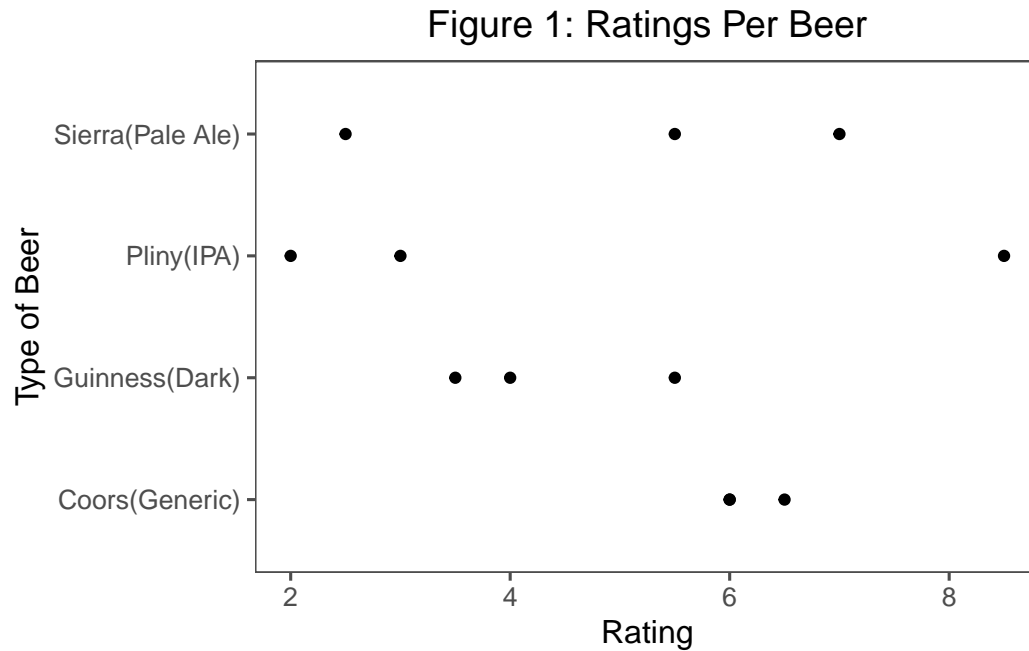


Figure 2 displays a visual representation of Ratings Per Subject. One possible trend observed in this chart is the different rating tendencies of the participants. Although the data overlaps, a potential trend could be that Beni naturally rates beers higher while Zach naturally rates them lower.




Table 4 displays the ANOVA results from the BIBD design. The p-value for the treatment (type of beer) is 0.40. and the p-value for the block (participant) is 0.34. Meaning that our factor of interest and

blocking factor are both insignificant.

Table 4: ANOVA Table for Linear Model

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Participant	3	12.17	4.06	1.20	0.40
Beer	3	14.40	4.80	1.42	0.34
Residuals	5	16.94	3.39	NA	NA

Post Hoc Analysis

Even though the ANOVA failed to reject our null hypothesis that all beer types were liked equally, it is still interesting to see confidence intervals for pairwise comparisons. We will use Tukey's HSD test, because we desire to look at all pairs of treatment means.

Table 5: CI Pairwise Comparison of Beers found with Tukey's HSD Test with sig. level = 0.05

Pairs	Estimated Difference	Lower Bound	Upper Bound
Guinness(Dark)-Coors(Generic)	-2.78	-8.32	2.77
Pliny(IPA)-Coors(Generic)	-2.17	-7.71	3.38
Sierra(Pale Ale)-Coors(Generic)	-1.72	-7.27	3.82
Pliny(IPA)-Guinness(Dark)	0.61	-4.93	6.16
Sierra(Pale Ale)-Guinness(Dark)	1.06	-4.49	6.60
Sierra(Pale Ale)-Pliny(IPA)	0.44	-5.10	5.99

Table 5 confirms our ANOVA results. That is, every confidence interval contains the value of zero. Interestingly, pairs involving Coors Light have the highest estimated difference. We recognize that the block factor is not the factor of interest here, but for the sake of the post hoc we'll calculate the CI's using Tukey's HSD as well.

Table 6: CI Pairwise Comparison of Participants found with Tukey's HSD Test with sig. level = 0.05

Pairs	Estimated Difference	Lower Bound	Upper Bound
Jon-Beni	-1.50	-7.05	4.05
Nolan-Beni	-1.67	-7.21	3.88
Zach-Beni	-2.83	-8.38	2.71
Nolan-Jon	-0.17	-5.71	5.38
Zach-Jon	-1.33	-6.88	4.21
Zach-Nolan	-1.17	-6.71	4.38

Table 6 shows us that there is no block effect. Looking at the estimated differences it seems that Nolan-Jon have the most similar taste, while Zach-Beni differ most out of the pairs. Of course, neither of these are statistically significant.

Conclusion

Main Results: The primary aim of this study was to identify the most favored beer type among a group of friends during Thursday Night Football watch parties using a Balanced Incomplete Block Design (BIBD). The results indicate that among the four types of beer tested—Guinness (Dark), Sierra Nevada (Pale Ale), Pliny the Elder (IPA), and Coors (Generic)—Coors received the highest mean rating (6.17), while Guinness received the lowest (4.33). However, the statistical analysis via ANOVA did not show significant differences in preference among the beer types (p -value = 0.40), nor did it show a significant effect from the participants as blocks (p -value = 0.34).

Most Important Results: The key takeaway from our experiment is that there was no statistically significant preference for any specific beer type among the participants. This suggests that within the small group of friends, beer preference might be more similar than anticipated, or perhaps the design did not capture the variability in taste preference. The lack of significance in the blocking effect indicates that individual participant biases did not significantly influence the ratings, suggesting that we might not need to block by participant in future experiments with similar designs.

Suggestions for Future Experiments: Increase Sample Size: The small number of participants (four) might not provide enough statistical power to detect differences reliably. A larger sample could help in distinguishing preferences more clearly. Expand Beer Selection: Including more varieties or brands within each beer category might reveal more subtle preferences among the participants. Better Control Palate Cleansing: Although water was provided to cleanse the palate, the sensory overload from tasting multiple beers might have affected later judgments. The use of bland crackers or taking breaks between tastings could better ensure an even judgement of each beer type. Power: For future experiments, the data obtained from this experiment could be used to help calculate the variance of the data which is required in power calculations.

Surprises: An unexpected result was the lack of a significant block effect. Given that individuals often have strong personal tastes in beer, we anticipated more variance due to participant differences. This might suggest that the beers were too similar in taste or that the participants were simply unable to detect significant differences in taste between the beers. The relatively small range in mean ratings (1.84 for beers, 2.83 for participants) was somewhat surprising, indicating perhaps a homogeneity in taste preference. This makes sense since all participants are of the same friend group and are very similar in age.

Data Concerns: Noise and Outliers: There were notable outliers in the ratings, such as scores of 8.5 and 2, which could skew results. However, these did not lead to significant outcomes due to the non-significant ANOVA results. This could imply either that the data was too noisy to discern clear preferences or that the variability in ratings was not substantial enough to be meaningful. Data Set Size: With only four participants, the dataset was inherently noisy and potentially less representative of broader beer preferences. This constraint might limit how broadly we can apply the study's conclusions.

In conclusion, while this experiment did not find significant differences in beer preferences, it has provided valuable insights into the application of BIBD in a social setting. Future research should consider these points to refine the study design for more precise and generalizable results.

Appendix

Code Used

Libraries Used

```
library(tidyverse)
library(ggthemes)
library(tidyr)
library(knitr)
```

Randomization Code

```
# Computing the BIBD matrix
trts <- c("A", "B", "C", "D")
set.seed(530)
t(replicate(4, sample(trts, 3, replace = FALSE)))

# Randomly matching treatments to the letters
brands <- c("Pliny", "Coors", "Siera", "Guinnes")
shuffled_brands <- sample(brands)
assignments <- data.frame(trts, shuffled_brands)

#Randomizing the order of participants/blocks
boys <- c("Zach", "Jon", "Nolan", "Benni")
rank <- 1:4
shuffled_names <- sample(boys)
order <- data.frame(rank, shuffled_names)
```

Data Code

```
# Data input
beers <- c("Sierra(Pale_ale)", "Coors(Generic)", "Guinness(Dark)", "Pliny(IPA)")
Nolan <- c(NA, 6, 5.5, 3)
Jon <- c(5.5, 6, 3.5, NA)
Benni <- c(7, NA, 4, 8.5)
Zach <- c(2.5, 6.5, NA, 2)

raw_data <- data.frame(beers, Nolan, Jon, Benni, Zach)
```

```
# Data Cleaning
pivoted_raw_data <- pivot_longer(raw_data,
                                  cols=-beers,
                                  names_to = "names",
                                  values_to = "rating")

cleaned_data <- pivoted_raw_data %>%
  drop_na(rating)
```

Plots

```
cleaned_data %>%
  ggplot(aes(x = Rating,
             y = Beer))+
  geom_point(size = 1.5)+
  theme_few()+
  ggtitle("Figure 1: Ratings Per Beer")+
  ylab("Type of Beer")+
  xlab("Rating")+
  theme(plot.title = element_text(hjust = 0.5))
```

```
cleaned_data %>%
  ggplot(aes(x = Rating,
             y = Participant))+
  geom_point(size = 1.5)+
  theme_few()+
  ggtitle("Figure 2: Rating Per Subject")+
  ylab("Participant")+
  xlab("Rating")+
  theme(plot.title = element_text(hjust = 0.5))
```

Anova Table

```
linear_model <- lm(Rating~Participant+Beer, data = cleaned_data)
anova_table <- anova(linear_model)
rounded_anova_table <- anova_table
kable(rounded_anova_table,
      caption = "ANOVA Table for Linear Model",
      digits = 2)
```

Post Hoc

```
# Tukey HSD analysis
model <- aov(Rating~Participant+Beer, data = cleaned_data)
test <- TukeyHSD(model)
```

Tables

```
#BIBD set (Not using the randomization code because
#set.seed() changes with different packages loaded)
x <- data.frame(p = c("1","2","3","4"),
               a = c("C", "B", "D", "D"),
               b = c("B", "C", "A", "B"),
               c = c("D", "A", "C", "A"))

kable(x,
      caption = "Blanced Incomplete Block Design Setup",
      col.names = c("Participant","I", "II", "III"),
      align = "c")

#Randomization of trts and blocks
y <- data.frame(Treatment = c("A", "B", "C", "D"),
               Beer = c("Sierra", "Coors", "Guinness", "Pliny"),
               Order = 1:4,
               Participant = c("Zach", "Jon", "Nolan", "Beni"))

kable(y,
      caption = "Randomization of Treatments and Participants") %>%
  column_spec(3, border_left = TRUE)

#Results
col_means <- round(colMeans(raw_data[, -1], na.rm = TRUE), 2)

raw_data_with_col_means <- rbind(raw_data, c("Mean", col_means))

row_means <- round(apply(raw_data[, -1], 1, mean, na.rm = TRUE), 2)

raw_data_with_means <- cbind(raw_data_with_col_means, Row_Mean = c(row_means, NA))

raw_data_with_means %>%
  kable(
    caption = "Beer Ratings by Participant",
    col.names = c("Beers", "Nolan", "Jon", "Beni", "Zach", "Row Means"),
```

```

    align = "c"
  )

#Anova Table
kable(rounded_anova_table,
      caption = "ANOVA Table for Linear Model",
      digits = 2)

#Tukey HSD tables
kable(test$Beer[1:6, 1:3],
      col.names = c("Pairs", "Estimated Difference", "Lower Bound", "Upper Bound"),
      align = "c",
      caption = "CI Pairwise Comparison of Beers
found with Tukey's HSD Test with sig. level = 0.05",
      digits = 2,
      escape = TRUE)
kable(test$Participant[1:6, 1:3],
      col.names = c("Pairs", "Estimated Difference", "Lower Bound", "Upper Bound"),
      align = "c",
      caption = "CI Pairwise Comparison of Participants
found with Tukey's HSD Test with sig. level = 0.05",
      digits = 2,
      escape = TRUE)

```